

Clustering Methods and Correlated Data: Simulations

Parker Knight

9/29/2021

Overview

1. Review of methods
2. Simulation design
3. Results
4. Directions

Methods of Interest

1. K-means
2. Hierarchical clustering
3. Spectral clustering

K-means

Let $X \in \mathbb{R}^{n \times p}$ be a given data matrix. Our goal is to partition the n data points into k non-overlapping clusters, denoted by the set $S = \{S_1, S_2, \dots, S_k\}$.

The k-means algorithm does so by minimizing the Euclidean distance between each point and the centroid of its assigned cluster:

$$\min_S \sum_{i=1}^k \sum_{X_j \in S_i} \|X_j - \mu_i\|_2^2$$

where μ_i is centroid of the points in S_i .

K-means (cont.)

Key points:

- ▶ k must be specified up front
- ▶ Computation is NP-hard (objective function is not convex)
- ▶ By default, uses Euclidean distances
 - ▶ alternate versions exist (K-medoids)
- ▶ Assumes separation between clusters is convex

Hierarchical clustering

Let $D = \text{dist}(X)$ be the pairwise distance matrix of the n samples.

In agglomerative HC, we start with every observation in a singleton cluster, join clusters G and H together based on the following criteria:

- i. $d_{SL} = \min_{i \in G, j \in H} d_{i,j}$ (single linkage)
- ii. $d_{CL} = \max_{i \in G, j \in H} d_{i,j}$ (complete linkage)
- iii. $d_{GA} = \frac{1}{N_G N_H} \sum_{i \in G} \sum_{j \in H} d_{i,j}$ (group average)

In divisive HC, we begin with data all in one cluster and split into sub-clusters with K-means.

HC (cont.)

Key points:

- ▶ don't need to specify k up front, can 'cut' the tree at any depth
- ▶ relies on pairwise distances, sensitive to outliers
- ▶ greedy algorithm
- ▶ 'big data' issues

Spectral clustering

From our data matrix X , compute a similarity matrix $S \in \mathbb{R}^{n \times n}$.
Then:

1. Compute a similarity graph from S , and define W to be the 'weighted' adjacency matrix
2. Construct the Laplacian L from W
3. Compute $L = UDU^T$
4. Cluster the first k columns of U using k-means

Spectral (cont.)

Key points:

- ▶ How do we construct S ? How do we construct a graph from S ?
- ▶ Normalized or unnormalized Laplacian?
- ▶ Need to specify k
- ▶ Can handle non-convexity between clusters

Simulation Design

Question: how robust are these methods to correlation structure between the samples in X ?

We vary the following parameters:

- ▶ $n = 100$
- ▶ $p = 50, 100$
- ▶ # of clusters = 2, 3, 4
- ▶ # of groups = 2, 10, 50
 - ▶ the 'groups' determine the block-diagonal structure of the covariance matrix
- ▶ $\rho = 0.2, 0.5, 0.8$
 - ▶ within-group correlation
- ▶ difference in mean between groups: 0.25, 1
 - ▶ 'signal strength'

Simulation Design (cont.)

Given a set of parameters, repeat R times:

1. Generate block-diagonal Σ with `ngroups` blocks, set within-group correlation ρ

```
Sigma <- Matrix::bdiag(lapply(1:(n/ngroups), function(i){  
  I <- (1-rho) * diag(ngroups)  
  ones <- rho * rep(1, ngroups) %*% t(rep(1, ngroups))  
  I + ones  
}))
```

Simulation Design (cont.)

2. Generate vector of cluster assignments $\mathbf{C} = (C_1, C_2, \dots, C_n)^T$ where C_i may equal $1, 2, \dots, k$. Set the mean vector $\mu = (\mu_{C_1}, \mu_{C_2}, \dots, \mu_{C_n})^T$ where μ_{C_i} is the mean of cluster C_i .

[illegible]

Simulation Design (cont.)

3. Draw $X \sim MVN(\mu, \Sigma)$.

```
X <- t(MASS::mvrnorm(n = p, mu = mu_vec, n/nclust,  
  Sigma = Sigma))
```

4. Cluster the rows of X with each method to get an estimated cluster assignment $\hat{\mathbf{C}} = (\hat{C}_1, \dots, \hat{C}_n)^T$ and compare results...

Simulation Design (cont.)

How do we evaluate the performance of each clustering method?

Rand index:

$$RI = \frac{1}{\binom{n}{2}} \sum_{i=1}^n \sum_{j \neq i} I \left[(C_i = C_j \wedge \hat{C}_i = \hat{C}_j) \vee (C_i \neq C_j \wedge \hat{C}_i \neq \hat{C}_j) \right]$$

Simulation Design (cont.)


Some immediate concerns:

1. We set k to be the true number of clusters
2. We assume that cluster means are equally spaced apart (signals are equally strong).
3. Many options for spectral clustering. . . our method (after some trial-and-error) is:
 - ▶ Construct S with a Gaussian similarity kernel with $\sigma^2 = 25$
 - ▶ Build an unweighted K-nearest neighbors graph with $K = 25$ (relatively dense)
 - ▶ Use a normalized Laplacian $L = I - D^{-1/2}AD^{-1/2}$


Results

clustering_method	average_ri
hclust	0.7221509
kmeans	0.7546772
spectral	0.6649400


Results (cont.)



Results (cont.)



Results (cont.)



Directions

1. Improve simulations
 - ▶ Bigger n and p
 - ▶ Move beyond Gaussian data
 - ▶ More complex correlation structures
2. Explore variations of spectral clustering
3. More clustering methods
 - ▶ Convex clustering?